

Brain Before the Swipe: Can Consumer-Grade EEG Predict Viewer Engagement during Naturalistic TikTok Browsing?

Master's Thesis

Gregor Lederer

Supervisors:

Dr. Irina-Catrinel Craciun

Prof. Dr. Dr. h. c. mult. Lothar H. Wieler

Prof. Dr. Bert Arnrich

A thesis presented for the degree of
Master of Science

Hasso Plattner Institute
University of Potsdam

Acknowledgements

I would like to express my sincere gratitude to my supervisors, and everyone else who supported me during this thesis.

Contents

Acknowledgements	I
Abstract	1
1 Introduction	3
2 Materials and Methods	4
2.1 Participants	4
2.2 Data Acquisition	4
2.3 Experimental Protocol	5
2.4 Labeling Strategy: The STAY Paradigm	7
2.4.1 Theoretical Justification for the STAY Objective	8
2.5 Internal Structuring: Python Pipeline Walkthrough	8
2.5.1 1. Signal Preprocessing & Normalization (<code>utils.py</code>) .	8
2.5.2 2. Standard Benchmark: Engagement Index (<code>step4_engagement_index.py</code>)	9
2.5.3 3. Feature Engineering: The 112 Dimensions (<code>step2_train_rf.py</code>)	10
2.5.4 4. Traversing Windows & Probabilistic Scaling (<code>step2_train_rf.py</code>)	10
2.5.5 5. Model Assembly & Validation Constraints	11

Abstract

This thesis investigates whether consumer-grade Electroencephalography (EEG) can accurately predict when a user is actively sustaining cognitive engagement during naturalistic TikTok browsing. Using a Muse S headband (4 channels, 256 Hz), we recorded brain activity from participants while they freely browsed their personal TikTok feed without artificial constraints. Behavior was mapped via tracking video transitions (screen swipes), allowing us to construct a rigorous binary classification framework focused on the underlying cognitive state: *sustained engagement* (STAY, Class 1) versus *imminent task abandonment* (SKIP, Class 0), captured via the 3-second pre-skip window.

We implemented a robust pipeline applying baseline-normalized relative spectral power features. We compared isolated feature representations across three machine learning architectures—a standard composite Engagement Index evaluated via Logistic Regression, a Baseline Decision Tree, and a Random Forest ensemble—evaluated in an intra-participant framework alongside cross-participant generalizability assessment via Leave-One-Group-Out cross-validation (LOGO-CV). The Random Forest (200 trees, 112 frequency-domain features) achieved the most robust performance, successfully diagnosing the STAY state with a high relative Sustained Engagement Recall intra-subject, significantly surpassing theoretical chance baselines.

Feature importance mapping revealed that localized high-gamma activity (40–60 Hz) at frontal (AF7/AF8) and temporal (TP9/TP10) nodes was universally the dominant predictive signature differentiating sustained viewing from an impending skip. An empirical ablation of the 50 Hz notch filter confirmed that this frequency range reflected genuine neurological state variance rather than power line contamination. The standard clinical Engagement Index ($\beta/(\alpha + \theta)$) fundamentally failed to discriminate these states natively, indicating that full structural spectral representations incorporating high-frequency volatility are required for micro-decision boundaries [?].

Behavioral skip-sequence analysis exposed target class heterogeneity, establishing that the predictive neural signature fundamentally contrasts gen-

uine content-locked sustained appraisal against a generalized “disengaged browsing” attentional mode. Consequently, these findings prove consumer-grade EEG can reliably identify individualized cognitive engagement states during naturalistic algorithmic media consumption, though cross-participant generalizability without per-user calibration remains fundamentally limited.

Chapter 1

Introduction

Short-form video platforms like TikTok now fundamentally shape how billions of people allocate their attention, often in bursts lasting mere seconds. On these platforms, every single swipe represents a critical micro-decision: to stay and engage, or to skip and move on. Repeated hundreds of times per session, these sequential micro-decisions determine exactly what information people absorb, which health or educational messages effectively reach them, and what digital behaviors are continuously reinforced.

For public health communication, this dynamic poses a significant challenge. Health promotion interventions delivered via social media must compete head-to-head with highly optimized entertainment content for the exact same limited cognitive resources. If a health promotion video fails to maintain engagement in the critical first few seconds, the core message never lands. To improve these interventions, we need to understand what actually causes a human brain to sustain engagement—not just what individuals retrospectively claim they liked, but rather what their objective neural state looked like during active, continuous viewing.

Understanding the neural signature of sustained engagement—and the exact moment that engagement begins to collapse—opens the door to designing communications that maintain attention precisely when the brain is most vulnerable to disengagement. This investigation is not about building manipulative technology, but rather about understanding how to make health-relevant content as neurologically compelling as the algorithms it competes against.

Chapter 2

Materials and Methods

2.1 Participants

Twenty-eight ($N = 28$) participants were recruited for the study ([Insert N_F female, N_M male]). To participate, individuals were required to have normal or corrected-to-normal vision and no history of neurological disorders. Three participants were excluded *a priori*: two due to a high number of self-reported keypress errors (greater than 10) which compromised label integrity, and one due to an undisclosed neurological condition. The experimenter’s own data, although collected for pipeline validation, was strictly excluded from all reported results to eliminate experimenter bias. The final sample consisted of $n = 25$ participants (age range 20–30 years, $M = 24.6$, $SD = 2.9$).

Participants were collected over a total of five days and were conceptually divided into two groups based on recruitment timing: an initial smaller group was recorded in a separate room without financial compensation, while the subsequent main group (recorded a few days later) received a 20-euro voucher as an incentive. All participants provided informed consent and signed data protection agreements approved by the Ethics Committee of the University of Potsdam prior to the study.

2.2 Data Acquisition

Electroencephalography (EEG) data were recorded using a consumer-grade Muse S headband (Generation 2). The device features four dry electrodes placed at positions AF7, AF8, TP9, and TP10 according to the international 10–20 system, with the reference electrode located at Fpz. Data were continuously sampled at a rate of 256 Hz.

The EEG signals were transmitted via Bluetooth and recorded using

[Insert Picture of Muse Headband]

Figure 2.1: The consumer-grade Muse S (Gen 2) EEG headband used for data acquisition.

the Lab Streaming Layer (LSL) protocol with the `muselsl` library. A custom Python recording script was developed to handle data collection. This script featured real-time frequency monitoring and an automatic reconnection mechanism to seamlessly manage any Bluetooth drops without disrupting the session.

To synchronize the behavioral data (video transitions) with the continuous EEG stream without relying on video monitoring of the participant, keystroke markers were mapped directly into the recording script. Participants used their non-dominant hand to press a specific key on a laptop keyboard provided in the room: the “B” key to denote the start of the baseline period, and the “A” key to mark every manual swipe (skip event) during the TikTok browsing session.

[Insert Picture of Experimental Setup Cabin]

Figure 2.2: The experimental setup cabin ensuring privacy and a distraction-free environment for the participant.

[Insert Closer Picture of Laptop and Phone Setup]

Figure 2.3: Close-up of the recording setup showing the researcher’s phone and the laptop used for keypress synchronization.

2.3 Experimental Protocol

The experimental sessions were scheduled for 45-minute timeslots, with the actual experimental procedures taking approximately 35 to 40 minutes. Upon arrival, participants were greeted, provided with an information sheet, and asked to sign the ethics consent forms. Afterward, they completed a pre-experiment demographic and basic screening survey, which took approximately five to ten minutes.

Participants were then instructed on the tasks. The Muse S headband was mounted on their head and adjusted to ensure optimal skin contact, particularly on the forehead and behind the ears for participants with long

[Insert Schematic Flowchart of Experimental Protocol]

Figure 2.4: Experimental protocol flow: Introduction and consent, EEG setup, 2-minute baseline video, task explanation, and 25-minute TikTok free browsing session.

hair. The experimenter verified the signal quality and ensured the device was adequately charged.

The experimental paradigm consisted of two main phases: a baseline recording and a free-browsing session. For the baseline phase, participants were provided with the researcher’s iPhone, and a two-minute nature video on YouTube was started. Participants were instructed to relax and simply watch the video. Concurrently, the recording script was initiated, and the “B” marker was pressed. To ensure ecological validity while keeping the participant focused, the iPhone’s “Guided Access” feature was enabled to prevent exiting the application. Participants were left alone in the room to afford them privacy. After two minutes, the Guided Access lock visually timed out, signaling the conclusion of the baseline period.

The experimenter then returned to explain the second phase: a naturalistic short-form video browsing task using the TikTok application. To ensure a standardized yet personalized starting state, the app’s “For You” feed was refreshed using the “Reset feed” option in the content preferences settings. Participants were instructed to consume TikTok content as they normally would, watching videos they found interesting and swiping away those they found boring. To restrict the action space strictly to scrolling behavior, they were asked to refrain from using the “like” function, writing comments, or visiting user profiles.

Crucially, participants were tasked with pressing the “A” button simultaneously with every swipe motion on the phone to accurately synchronize their micro-decisions with the EEG data. For instance, a participant holding the phone in their right hand used their left hand to press the designated key on the adjacent laptop keyboard.

[Insert Schematic of Button Click Synchronization]

Figure 2.5: Schematic representation of the manual synchronization process: the participant swipes to skip the video and simultaneously presses the ‘A’ button to mark the event in the EEG stream.

The Guided Access feature was reactivated and set to a 25-minute limit for the TikTok app. Participants were reminded that the headband script would automatically reconnect if a Bluetooth disconnection occurred. They

were then left alone in the room with the door closed, creating an isolated viewing environment designed to elicit natural responses to the video feed. They were instructed to knock on the glass door if any errors or issues arose; however, no significant interruptions occurred during the sessions.

Once the 25-minute viewing task concluded, the experimenter entered the room, and the participant removed the headband. They then answered a few post-session survey questions, specifically self-reporting if they made any keypress synchronization errors during the session. Finally, they signed the required compensation forms to receive their 20-euro voucher. The post-measurement debrief and surveying process took roughly two to four minutes. Participants were thanked for their time, briefly debriefed about the broader context of the study if interested, and dismissed.

2.4 Labeling Strategy: The STAY Paradigm

To translate the continuous EEG recordings into a supervised machine learning task, the data were segmented based on the keystroke markers. The primary goal of this research is not simply to predict the mechanical action of a thumb swiping, but rather to identify when a user is actively holding their attention on the content. Therefore, the computational target strictly isolates the underlying cognitive state into two discrete classes: the state of sustained engagement (**STAY**, designated mathematically as Class 1) and the state of impending disengagement (**SKIP**, designated mathematically as Class 0).

The **SKIP** class (Class 0) was strictly defined as the finite 3.0-second window terminating exactly at the timestamp of each registered A keypress (swipe event). This specific window captures the neural activity immediately preceding the physical action of skipping the video.

The **STAY** class (Class 1) represents sustained cognitive engagement. It was defined by extracting all other continuous viewing periods during the naturalistic TikTok consumption phases that fell entirely outside of the 3.0-second pre-skip boundaries.

The manually categorized pre-task resting baseline phases were structurally excluded from the machine learning training data. This guaranteed that the algorithms were explicitly trained to differentiate between active engagement and active disengagement while browsing, rather than simply distinguishing resting from viewing.

2.4.1 Theoretical Justification for the STAY Objective

The decision to mathematically map sustained engagement as the positive target class (Class 1) ensures that the machine learning models strictly evaluate our core behavioral science objective: measuring human attention capacity. Within this paradigm, the physical swipe is viewed merely as the delayed behavioral exhaust—the physical evidence that the cognitive state of active engagement has already critically failed [?].

The choice to isolate the **SKIP** sequence specifically to a 3.0-second anticipatory temporal window is firmly anchored in classical motor preparation and intentional neurological mapping literature [?]. Research reliably demonstrates that pre-cognitive physiological markers isolating imminent voluntary mechanical decisions manifest deep within cortical readiness potentials significantly ahead of downstream physical muscular executions. Consequently, isolating the exact 3.0 seconds preceding a mechanical swipe captures the moment of cognitive attention decay and motor execution readiness, decisively decoupling the mental decision to skip from the physical contamination (motor artifacts) generated by the arm movement itself.

2.5 Internal Structuring: Python Pipeline Walk-through

Note: The following subsections describe step-by-step exactly what the Python code in the `04-26_data_analysis_and_results` directory executes to process the EEG signals and build the machine learning models. This is for internal structural tracking and mathematical comprehension.

2.5.1 1. Signal Preprocessing & Normalization (`utils.py`)

- **True Sampling Rate Extraction:** The script does not assume the hardware delivers a perfect 256 Hz. Instead, it extracts the true, effective sampling rate by calculating the median time difference (Δt) between consecutive Lab Streaming Layer (LSL) timestamps:

$$f_{\text{effective}} = \frac{1}{\text{median}(\Delta t)}$$

- **Frequency Band Splitting:** The raw EEG wave is fed through a continuous 4th-order Butterworth bandpass filter. It systematically slices the wave into 7 established frequency bounds: Delta (1–4 Hz),

Theta (4–8 Hz), Alpha (8–13 Hz), Beta (13–30 Hz), Low Gamma (30–40 Hz), High Gamma (40–60 Hz), and Very High (60–100 Hz).

- **Notch Filter Exclusion:** While a 50 Hz power-line notch filter exists in the script, it is actively disabled via the `-nonotch` parameter to avoid systematically destroying the target High Gamma neurological signal.
- **Baseline Relative Power Normalization:** To cancel out physiological noise (like skull thickness or skin resistance unique to each participant), the pipeline uses the 100-s baseline:

1. **Absolute Power Conversion:** First, the script isolates the instantaneous voltage amplitude $A(t)$ and squares it to derive physical electrical power:

$$P_{abs}(t) = A(t)^2$$

2. **Resting Geometric Mean:** For all 28 extracted channels, it averages the absolute power specifically across the initial 100-second resting baseline ($\mu_{baseline}$). A failsafe sets the minimum bound to 1×10^{-12} to computationally prevent division by zero.
3. **Relative Scaling:** During the active TikTok viewing task, every data point is divided by this resting mean. This converts the raw waves into a clean, personalized "Relative Power Ratio":

$$P_{rel}(t) = \frac{P_{abs}(t)}{\mu_{baseline}}$$

2.5.2 2. Standard Benchmark: Engagement Index (step4_engagement_index)

- To test classic clinical theories against the EEG data, the script computes the standard Muse Engagement Index. For every 3.0-second window, it mathematically averages the relative power for the Alpha, Theta, and Beta bands across all four electrodes.
- The final scalar metric is computed by dividing the active Beta power by the sum of the passive Alpha and Theta power:

$$EI_{standard} = \frac{\bar{P}_{\beta}}{\bar{P}_{\alpha} + \bar{P}_{\theta}}$$

- This single EI value is passed structurally to an established `LogisticRegression` model using 5-fold Stratified Cross-Validation (on balanced `STAY/SKIP` data). This serves as the baseline diagnostic proof testing if simple classic formulas can accurately detect the `STAY` state.

2.5.3 3. Feature Engineering: The 112 Dimensions (`step2_train_rf.py`)

- **Statistical Compression:** Because raw time-series data is incredibly volatile, the Random Forest pipeline utilizes statistical extraction. Every 3.0-second window consists of roughly 768 physical sampling timestamps across 28 relative power channels ($4 \text{ electrodes} \times 7 \text{ bands}$).
- Instead of using all 768 points, the script condenses the 3.0-second sequence along four distinct descriptive parameters describing the geometry of the wave over that window:
 1. **Mean** (μ)
 2. **Standard Deviation** (σ)
 3. **Minimum Boundary** ($\min$)
 4. **Maximum Boundary** ($\max$)
- **Total Dimension Yield:** $7 \text{ physiological bands} \times 4 \text{ hardware channels} \times 4 \text{ statistical distributions} = \mathbf{112 \text{ continuous input features}}$.

2.5.4 4. Traversing Windows & Probabilistic Scaling (`step2_train_rf.py`)

- **Extracting the SKIP class (Class 0):** The script drops a strict 3.0-second window terminating sequentially at the exact moment of every A keypress (swipe). Meaning exactly ONE data block is extracted per skip operation.
- **Extracting the STAY class (Class 1):** To map sustained engagement, the script shifts a 3.0-second window through all the remaining continuous viewing periods. It operates with an aggressive **80% overlap**, shuffling the window forward by a tiny 0.6-second stride. This yields massive amounts of overlapping STAY arrays.
- **50/50 Algorithmic Rebalancing:** The overlapping logic inherently creates a wildly imbalanced dataset mapping thousands of STAY windows against a few dozen SKIP windows. Using `numpy.random.choice` (Seed = 42), the script systematically deletes random STAY windows until the population perfectly matches the SKIP window counts exactly, rigorously enforcing a 50/50 prior distribution to physically prevent classifier bias.

2.5.5 5. Model Assembly & Validation Constraints

- **Intra-Subject Consistency (`step2_train_rf.py`):**
 - For every single participant individually, the balanced 50/50 array is chopped into a 60/40 Train/Validation Split.
 - A `RandomForestClassifier` is deployed utilizing exactly 200 estimators (`n_estimators=200`), constrained at a `max_depth=7` and `min_samples_leaf=5` to aggressively prevent overfitting the high-dimensional data.
 - The model performance is validated natively using the Recall (Sensitivity) locked explicitly tracking the **STAY** (Class 1) target label.
- **Inter-Subject Generalizability via LOGO-CV (`step3_logo_cv_rf.py`):**
 - To test whether the predictive map scales blindly across humans, the script iterates through Leave-One-Group-Out Cross Validation (LOGO-CV).
 - In a 25-step loop, it tosses 24 individuals into a master training bucket, explicitly rebalances that massively mixed array back to 50/50 natively, trains the Random Forest, and attempts mathematically to diagnose the **STAY/SKIP** sequences of the hidden 25th participant sequentially.

Bibliography